










DFlash speculative decoding — LongBench-E benchmark-wide averages (Qwen3-8B vs Qwen3.5-9B @ 4K / 8K / 16K / 32K / 64K, block=16, gamma=15, greedy)



qwen3-8b · 4K qwen3-8b · 8K qwen3-8b · 16K qwen3-8b · 32K qwen3-8b · 64K
qwen3.5-9b · 4K qwen3.5-9b · 8K qwen3.5-9b · 16K qwen3.5-9b · 32K qwen3.5-9b · 64K
qwen3-8b (trend) qwen3.5-9b (trend)

Samples per context length [8B / 9B where they differ] — 4K: n=32 (13 datasets) · 8K: n=32 (12 datasets) · 16K: n=32 (10 / 9 datasets) · 32K: n=32 (9 datasets, 31 composed) · 64K: n=32 (8 datasets, 32 composed)
 Mean output tokens per request [8B / 9B] — 4K: 178 / 246 · 8K: 183 / 185 · 16K: 218 / 246 · 32K: 12 / 24 · 64K: 49 / 24 (long-context runs emit far fewer tokens; compare per-token metrics, not per-request ones)