



DFlash speculative decoding — LongBench-E benchmark-wide averages (Qwen3-8B vs Qwen3.5-9B @ 4K / 8K / 16K, block=16, gamma=15, greedy)

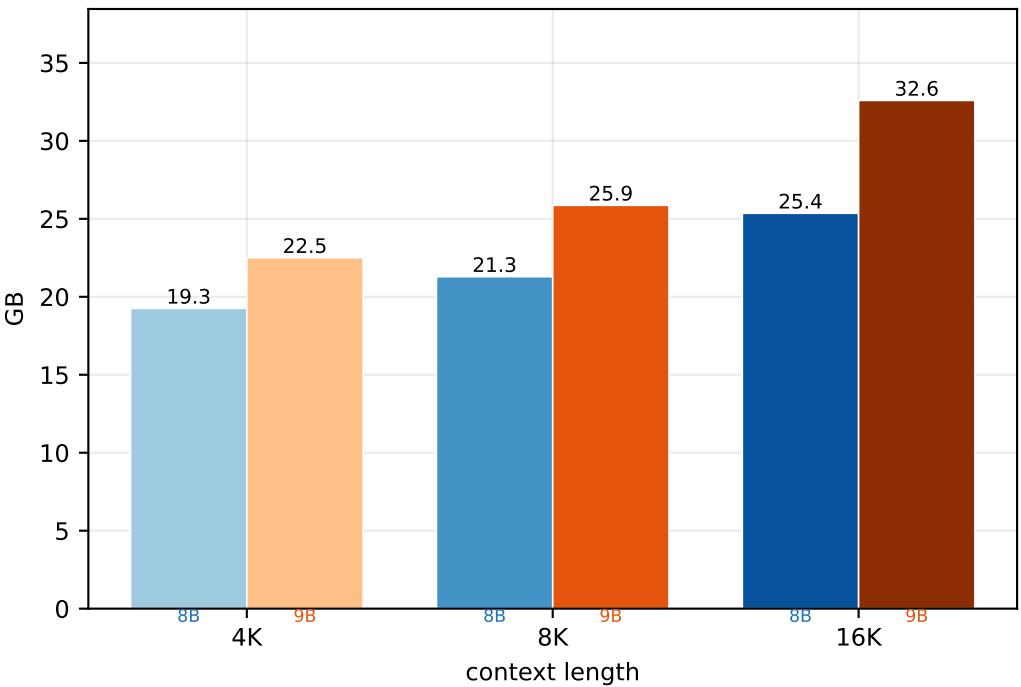
Mean acceptance length
(accepted tokens per verify step)



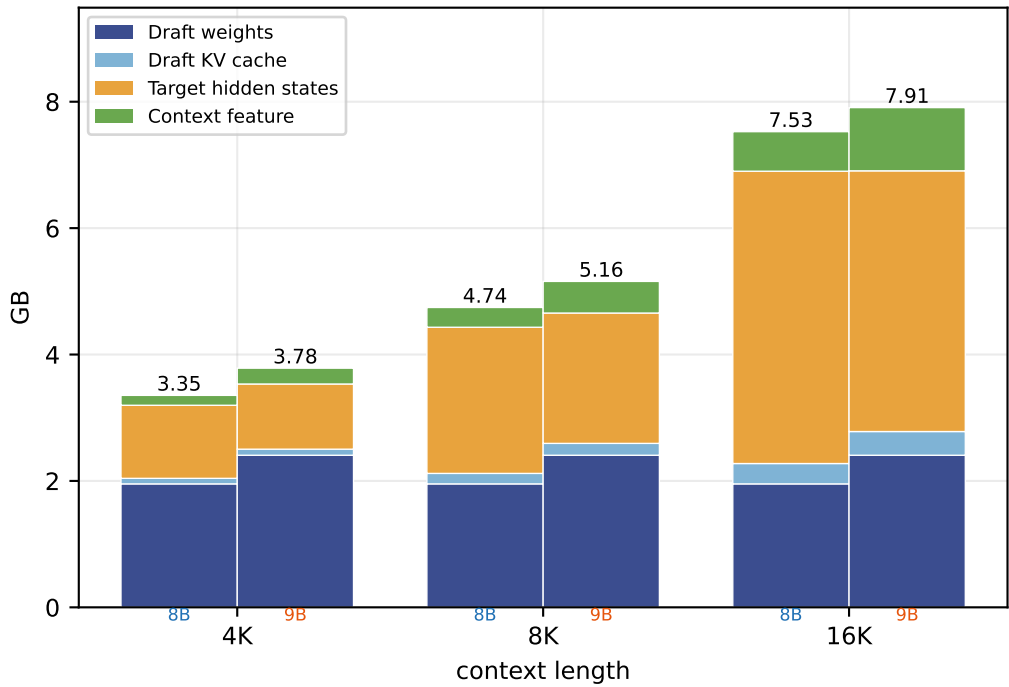
Acceptance length distribution
(16 = full γ block accepted)



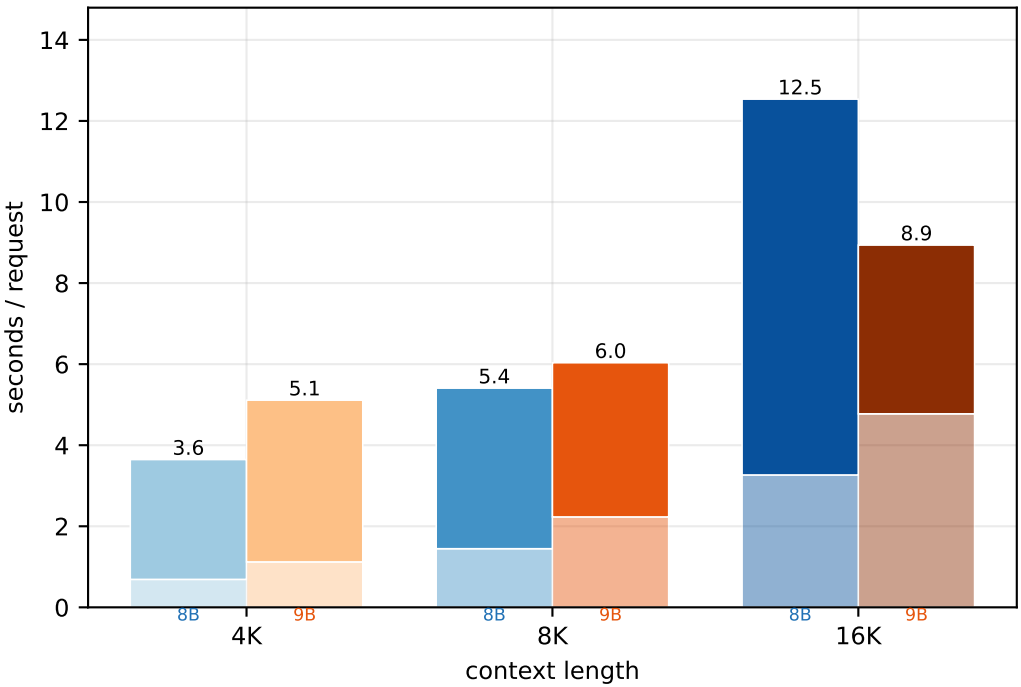
Peak GPU memory
(allocated, incl. draft overhead)



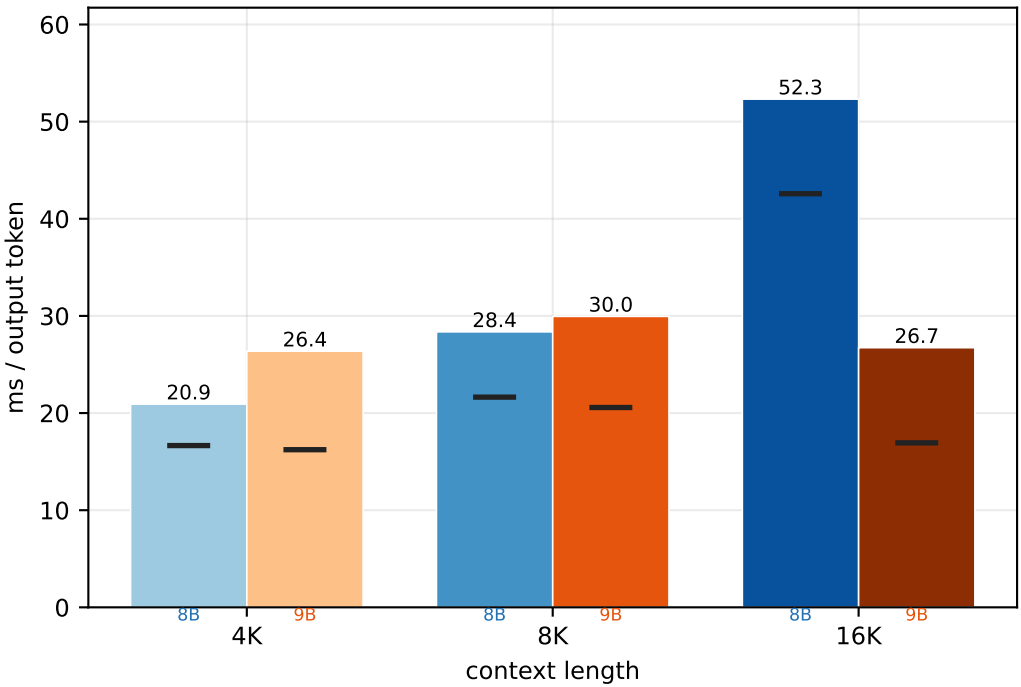
Draft-side memory overhead breakdown
(total = draft_overhead_gb)



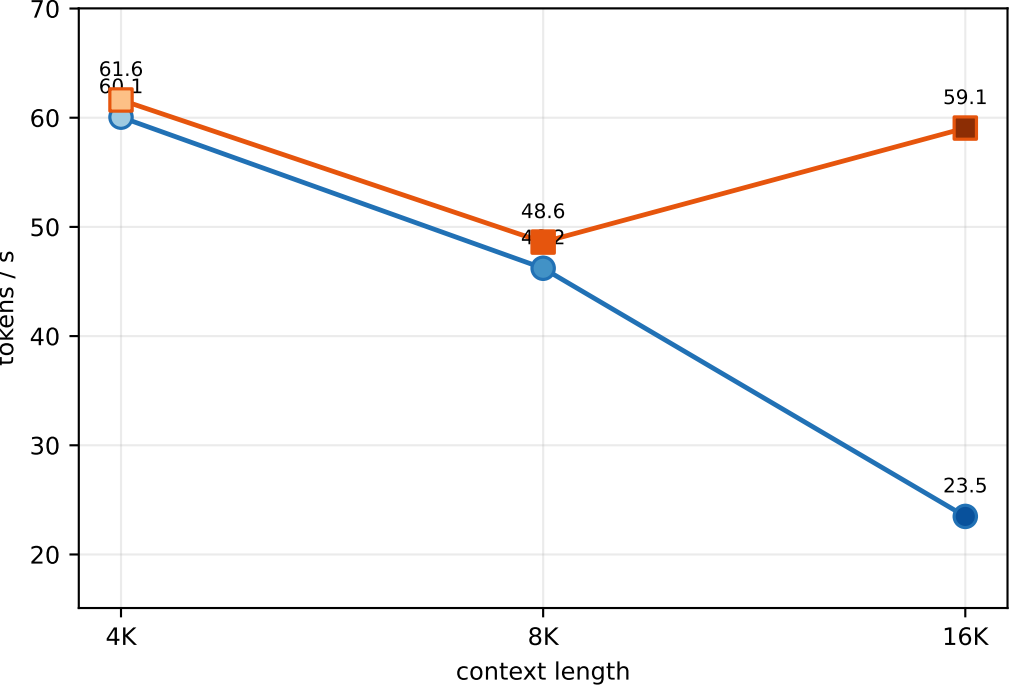
Mean end-to-end latency
(light = TTFT, solid = decode)



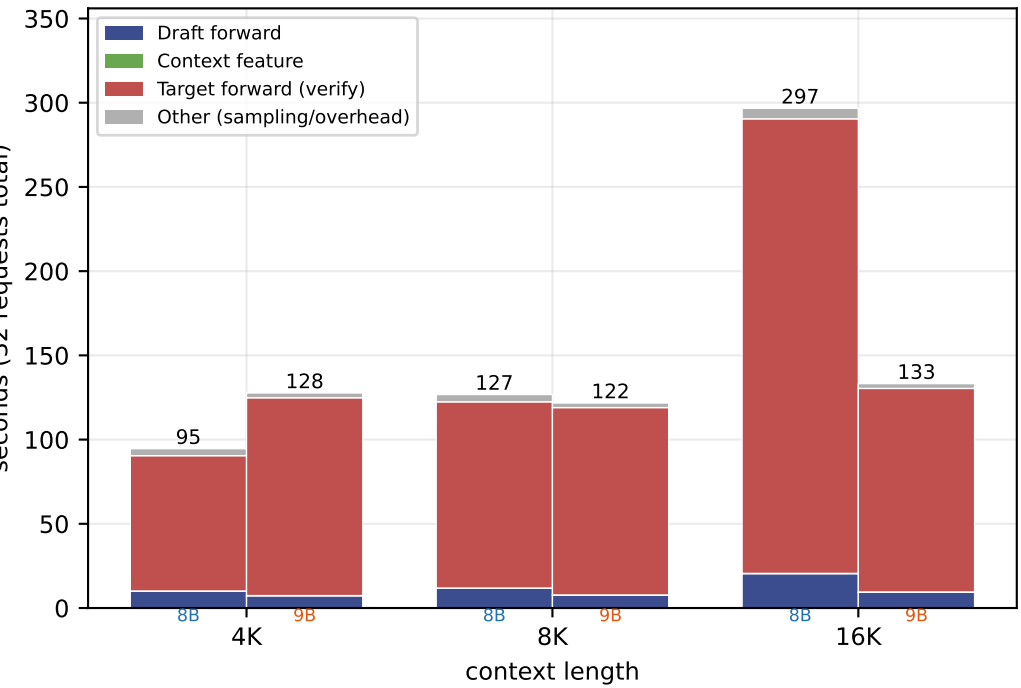
Per-token decode latency
(bar = per-request mean, — = aggregate over all tokens)



Decode throughput
(output tokens / s)



Decode wall-clock breakdown
(sum over 32 requests)



Drafter share of decode time
(draft forward + context feature)

