



DFlash vs. baseline — qwen3-8b on LongBench-E, benchmark-wide averages (4K / 8K / 16K / 32K / 64K; draft-only metrics excluded)


Mean end-to-end latency
(light = TTFT, solid = decode)




Mean TTFT
(prefill, identical work in both modes)




Per-token decode latency
(bar = per-request mean, — = aggregate)




Decode throughput
(output tokens / s)




Peak GPU memory
(bar = allocated, — = reserved)




Max target KV cache
(same target model, same context)



Target forward passes per request
(DF = verify steps, BL = decode steps)



DFlash relative to baseline



DFlash · 4K DFlash · 8K DFlash · 16K DFlash · 32K DFlash · 64K Baseline · 4K Baseline · 8K Baseline · 16K Baseline · 32K Baseline · 64K DFlash (trend) Baseline (trend)

Samples per context length (qwen3-8b) — 4K: n=32 (13 datasets) · 8K: n=32 (12 datasets) · 16K: n=32 (10 datasets) · 32K: n=32 (9 datasets, 31 composed) · 64K: n=32 (8 datasets, 32 composed)

Mean output tokens per request [DFlash / baseline] — 4K: 178 / 176 · 8K: 183 / 190 · 16K: 218 / 218 · 32K: 12 / 12 · 64K: 49 / 50 (long-context runs emit far fewer tokens; compare per-token metrics, not per-request ones)