


DFlash vs. baseline — qwen3-8b on LongBench-E, benchmark-wide averages (4K / 8K / 16K / 32K / 64K; draft-only metrics excluded)



Samples per context length (qwen3-8b) — 4K: n=32 (13 datasets) · 8K: n=32 (12 datasets) · 16K: n=32 (10 datasets) · 32K: n=32 (9 datasets, 31 composed) · 64K: n=32 (8 datasets, 32 composed)

Mean output tokens per request [DFlash / baseline] — 4K: 178 / 176 · 8K: 183 / 190 · 16K: 218 / 218 · 32K: 12 / 12 · 64K: 49 / 50 (long-context runs emit far fewer tokens; compare per-token metrics, not per-request ones)