


Qwen3.5-9B — verify forward: memory-bound → compute-bound as the batch grows (RTX A6000, N = 256, EOS suppressed)



Bytes and FLOPs are analytic (plan §5.4): executed target weights once (no embedding table), KV at the mean context $S + N/2$ once per row, GDN state in and out, logits written and read; FLOPs = 2-params-tokens + attention $4 \cdot q \cdot L \cdot H \cdot q \cdot d$ per layer + the delta rule. The draft adds fc and K/V projection of the kept context columns and 15 LM-head rows. Time is the CUDA-event mean per step, over all steps of the run (a finished row still computes). 9B ran on fla 0.5.2 kernels. A DFlash verify sitting below both roofs is the GDN block, whose per-layer kernels leave the card idle between launches — see batch_verify_breakdown, not the batch.