


Qwen3.5-9B — throughput, speedup, acceptance and the static-batch tail (32 prompts, N = 256 per row)



Makespan = all 32x256 tokens over the decode wall time; a static batch runs until its slowest row, and a DFlash row finishes after 256/ τ steps that vary by prompt, so the tail is DFlash's. 'All rows active' uses only the steps before the first row of a batch finished. τ counts the bonus token ($\tau = 1 + \text{accepted}$). Output agreement is exact greedy equality in bf16: baseline against itself at B=1 disagrees about as often as DFlash against baseline, so the gap is batch-shape numerics, not a lossless-ness bug (fp32 CPU tests match exactly). The prompt mix differs by S, so τ across S mixes length with task.