



Long-context memory: prompt recording dominates hybrid speculation state

● Qwen3-8B · full attention + dense ■ Qwen3.5-9B · GDN + dense ▲ Qwen3.5-35B-A3B · GDN + MoE

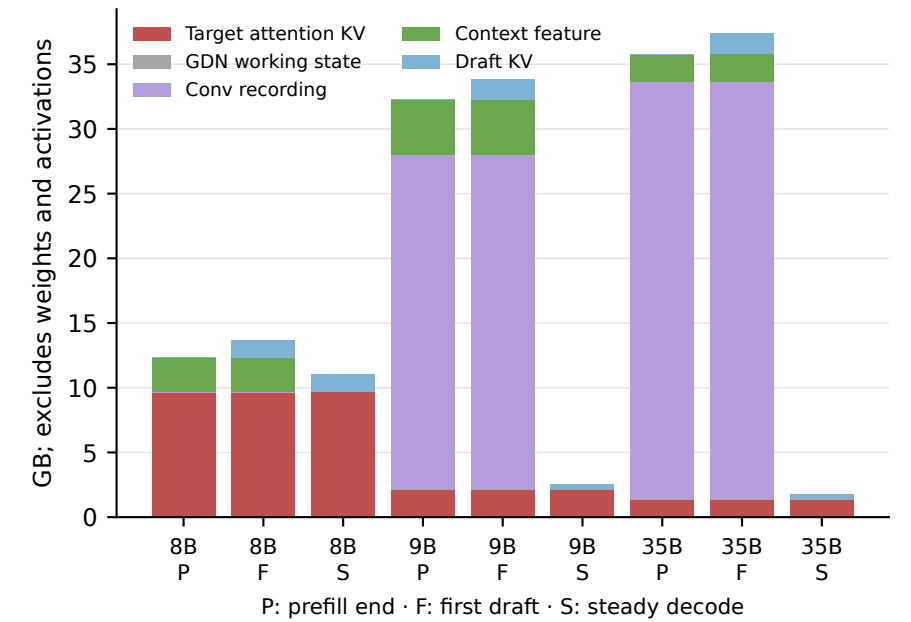
A Attention KV vs prompt conv recording



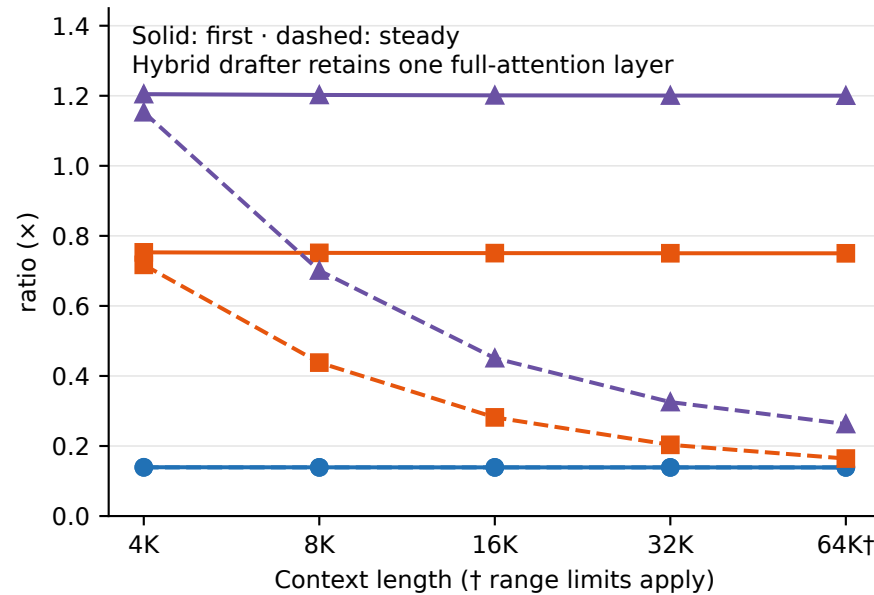
B Prompt context feature persists at first draft



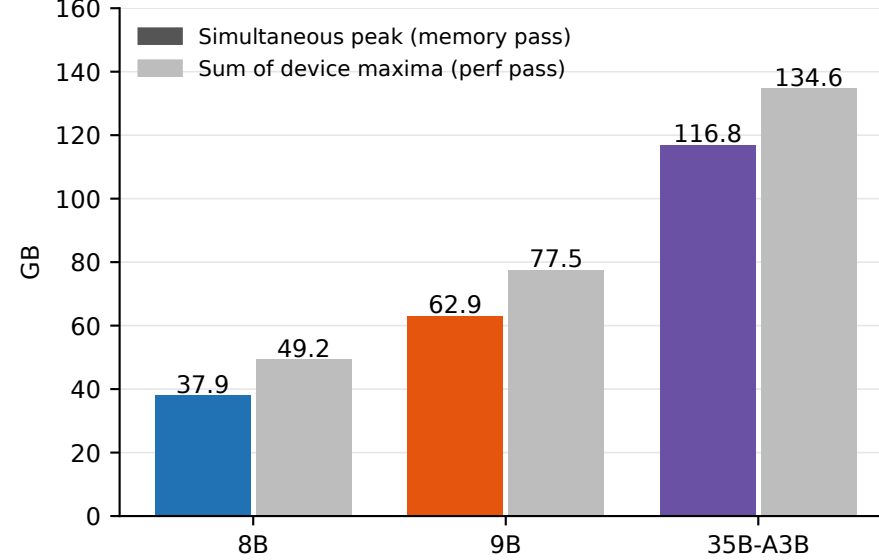
C 64K storage at three recorded moments



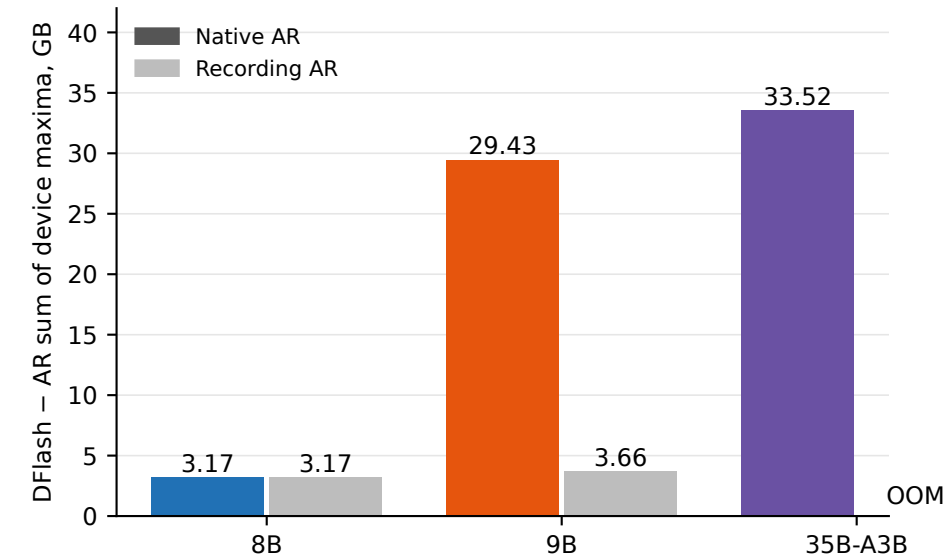
D Draft KV / target attention KV



E 64K peak definitions are different



F 64K overhead depends on AR policy



B=1 · 3 × RTX A6000, layer sharding · greedy · 256 output tokens · causal_conv1d+fla
One source document, three timing repeats. Hybrid results use stock rollback; lossless speedup is unverified.
64K: 8B exceeds configured RoPE range; 35B drafter exceeds stated training context. No compute-bound claim.