



Long-context DFlash: stock performance and the cost of verification

Qwen3-8B · full attention + dense Qwen3.5-9B · GDN + dense Qwen3.5-35B-A3B · GDN + MoE

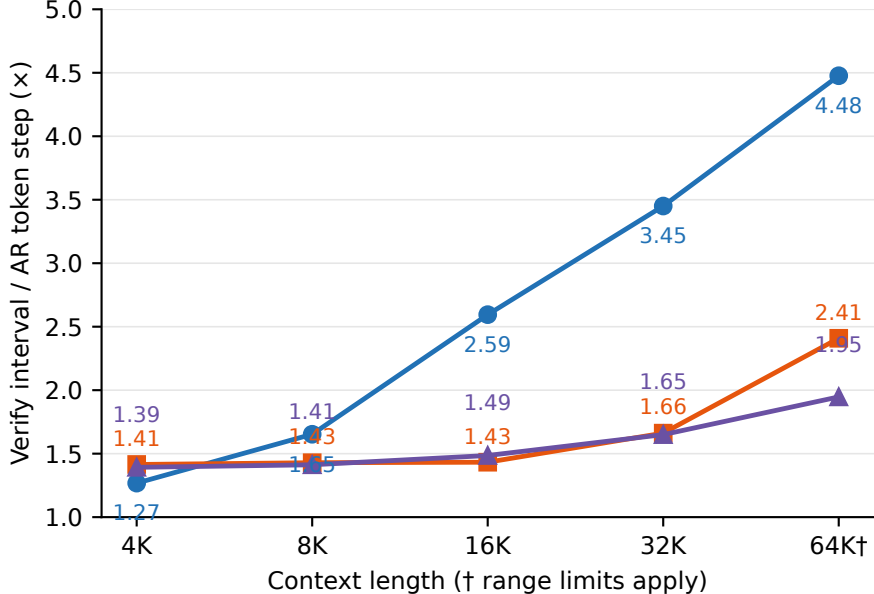
A Decode speedup vs native AR



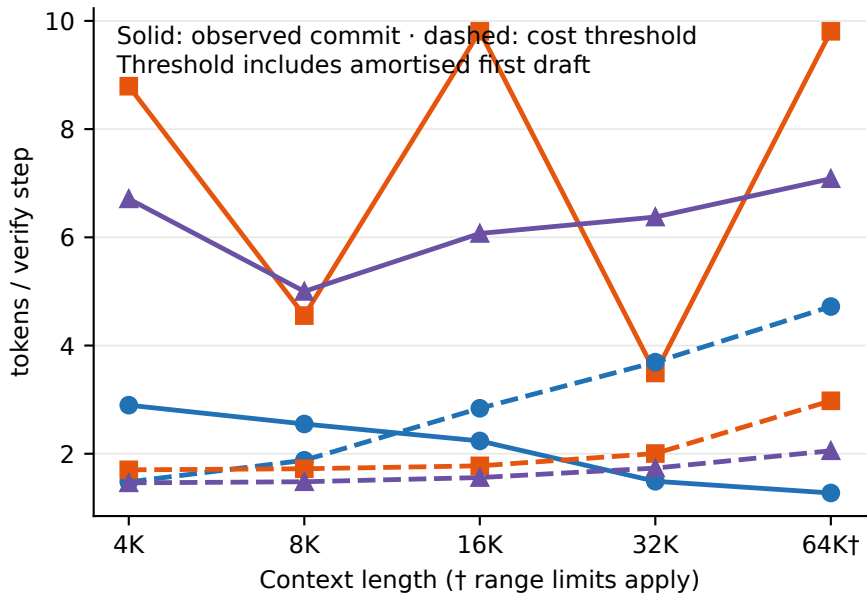
B Per-token latency (min-max over 3 repeats)



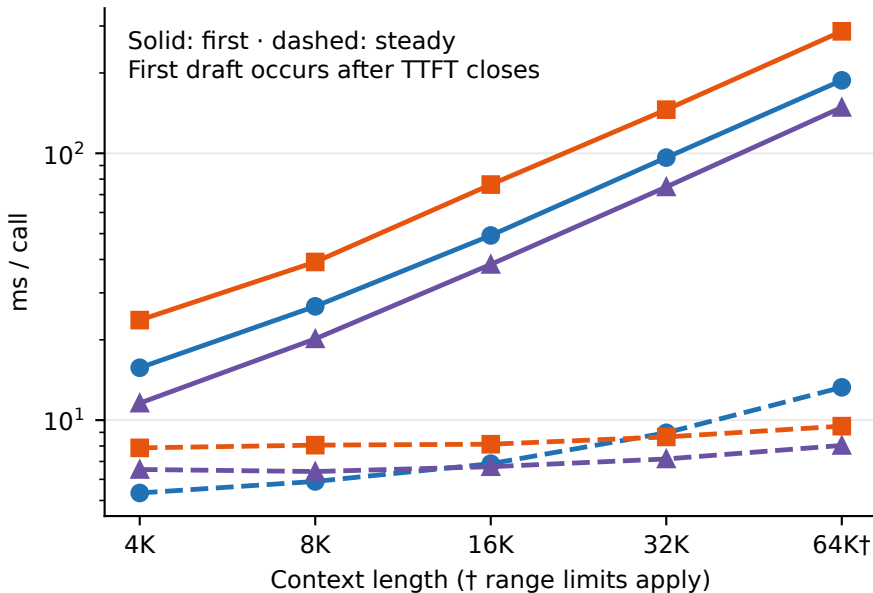
C Verify cost grows with context



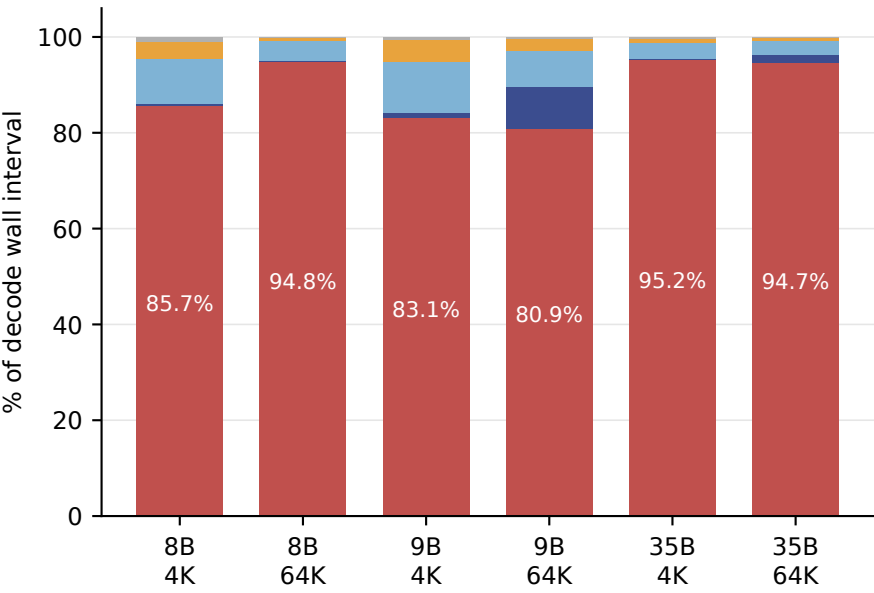
D Committed tokens vs break-even cost



E First and steady draft forwards



F Decode composition at 4K and 64K



B=1 · 3 × RTX A6000, layer sharding · greedy · 256 output tokens · causal_conv1d+fla
One source document, three timing repeats. Hybrid results use stock rollback; lossless speedup is unverified.
64K: 8B exceeds configured RoPE range; 35B drafter exceeds stated training context. No compute-bound claim.

Verify Steady draft Context feature
First draft Draft logits Remainder