



# Verification: residual attention growth and sparse expert execution

● Qwen3-8B · full attention + dense    ■ Qwen3.5-9B · GDN + dense    ▲ Qwen3.5-35B-A3B · GDN + MoE


**A 35B verify GPU kernel breakdown**




**B Full-attention mixer vs GDN mixer**




**C 35B GPU kernel sums and perf intervals**




**D 32K block sweep: latency per output token**



**E 35B: expert union broadens with block width**



**F Each hit expert still receives few tokens**



B=1 · 3 × RTX A6000, layer sharding · greedy · 256 output tokens · causal\_conv1d+fla  
One source document, three timing repeats. Hybrid results use stock rollback; lossless speedup is unverified.  
64K: 8B exceeds configured RoPE range; 35B drafter exceeds stated training context. No compute-bound claim.