










DFlash vs. baseline — qwen3.5-9b on LongBench-E, benchmark-wide averages, hidden_states=selective (4K / 8K / 16K / 32K / 64K; draft-only metrics excluded)



DFlash · 4K DFlash · 8K DFlash · 16K DFlash · 32K DFlash · 64K Baseline · 4K Baseline · 8K Baseline · 16K Baseline · 32K Baseline · 64K DFlash (trend) Baseline (trend)

Samples per context length (qwen3.5-9b) — 4K: n=32 (13 datasets) · 8K: n=32 (12 datasets) · 16K: n=32 (9 datasets) · 32K: n=32 (9 datasets, 31 composed) · 64K: n=32 (8 datasets, 32 composed)
Mean output tokens per request [DFlash / baseline] — 4K: 246 / 178 · 8K: 185 / 149 · 16K: 246 / 230 · 32K: 24 / 11 · 64K: 24 / 11 (long-context runs emit far fewer tokens; compare per-token metrics, not per-request ones)

Panels highlighted in green differ from visualization/ (hidden_states=full) — ADDED = a quantity that did not exist in those records, REVISED = same quantity, changed measurement

- [1] REVISED hidden_states=selective. The drafter reads the target's residual streams through forward hooks on the injected layers rather than output_hidden_states=True, so the tuple the target materialises drops from L_t+1 layers to n_inj. Only DFlash pays that term, so only the DFlash bars move; the baseline's peak is unchanged from the full-mode records. Acceptance and output tokens are identical in both modes.
- [2] ADDED summary*.memory_budget and target_weight_gb, new record fields, for both configurations. The baseline panel lists the draft weights it never calls, because the drafter is resident on the same card for the whole benchmark -- which is why the peak delta between the two bars is not the drafter's cost.
- [3] REVISED Device count differs from the full-mode records: 9B/32K 2->1, 9B/64K 3->2. Selective frees enough memory to fit on fewer cards. 9B/64K is still sharded and starred on the x axis, so its timings are pipeline-parallel. Memory stays comparable across a device change; timings do not -- going from sharded to single-GPU speeds the *baseline* up more than it speeds DFlash, so a speedup can read lower without DFlash having got slower.